Alibaba Group

December Quarter 2025 Results

Conference Call

Transcript

Thursday, 19 March 2026

Introduction

Lydia Liu

Head of Investor Relations, Alibaba Group

(Original)

Good day everyone and welcome to Alibaba Group's December Quarter 2025 earnings conference call. Joining us today are Joe Tsai, Chairman; Eddie Wu, Chief Executive Officer; Toby Xu, Chief Financial Officer; Jiang Fan, Chief Executive Officer of Alibaba E-Commerce Business Group.

I would like to remind you that this call is also being webcast on our corporate website. A replay of the call will be available on our website later today.

Now I will quickly cover the safe harbor. Today's discussions may contain forward-looking statements, particularly statements about our business prospects and expected financial results, that are subject to risks and uncertainties, which could cause actual results to differ materially from those contained in the forward-looking statements. Please refer to the safe harbor statements that appear in our press release and investor presentation provided today. Please note that certain financial measures that we use on this call are expressed on a non-GAAP basis. Our GAAP results and reconciliations of GAAP to non-GAAP measures can be found in our earnings press release and investor presentation provided today.

Unless otherwise stated, growth rate of all stated metrics mentioned during this call refers to year-over-year growth versus the same quarter last year

And now I will turn the call over to Eddie.

Group Business Update

Eddie Wu

CEO, Alibaba Group

(Translation)

Welcome, and thank you for joining our quarterly earnings call.

Over the past quarter, we continued to invest decisively in our two strategic priorities — AI+Cloud and consumption. Cloud Intelligence Group revenue growth accelerated to 36%, while our quick commerce business continued to expand in scale with ongoing improvement in unit economics.

With the dawn of the AI agent era in 2026, the addressable market for AI infrastructure providers like Alibaba is set to grow exponentially. AI models and capabilities are rapidly being embedded into all kinds of mainstream work environments, with token consumption surging across sectors. Historically, cloud and software budgets for enterprise IT services have represented only around 5% of corporate revenue. As model-driven agents begin to handle mainstream tasks across industries, our total addressable market will expand severalfold.

From AI infrastructure to models and on to the application layer, Alibaba has built a complete full-stack AI capability set to support this exponential growth in demand. Faced with an industry transformation and strategic opportunity of this magnitude, Alibaba Group is in a new phase of entrepreneurial reinvention and critical investment oriented toward the future.

Next, let me outline Alibaba's AI strategic roadmap.

We possess full-stack AI capabilities: chips and cloud computing form the infrastructure layer, while the AI model and application layer — anchored by Alibaba Token Hub — comprises foundation models, MaaS, and both enterprise and consumer applications. Together, these give us end-to-end coverage and use cases across the entire stack, from AI infrastructure to models and applications.

Given the long-term growth momentum of the AI market, combined with Alibaba's full-stack positioning across the AI value chain, our business goal is clear: over the next 5 years, we aim to surpass US$100 billion in combined cloud and AI external commercial revenue, including MaaS.

At the infrastructure layer, driven by sustained strong AI demand, Cloud Intelligence Group's revenue from external customers accelerated to 35% growth this quarter. AI-related product revenue delivered triple-digit year-over-year growth for the tenth consecutive quarter. Cloud Intelligence Group's market share has grown for three consecutive quarters, reaching 36% and further widening our lead. Alibaba Cloud's cumulative external commercial revenue in fiscal year 2026 through to the end of February officially surpassed RMB 100 billion.

Over the past three months, average daily token consumption for public model APIs on our Model Studio (Bailian) MaaS platform grew sixfold from December to March. We expect commercial MaaS to become the largest revenue-generating product within Cloud Intelligence Group.

Our proprietary T-Head AI chips have achieved mass production at scale: as of February 2026, we had cumulatively shipped 470,000 chips. In real-world business deployments through Alibaba Cloud, more than 60% of T-Head chips serve external commercial customers, supporting AI workloads for over 400 enterprises across industries including internet, financial services, and autonomous driving. We are confident that T-Head's compute supply will continue to expand, contributing high-quality compute to our cloud infrastructure and MaaS platform, and strengthening the competitiveness of our cloud services.

At the model and application layer, we established the new Alibaba Token Hub (ATH) Business Group with the core mission of "creating, delivering, and applying tokens". ATH comprises Tongyi Laboratory, the MaaS Business Line, the Qwen Business Unit, the Wukong Business unit, and the AI Innovation Business Unit. This provides the organizational foundation for executing Alibaba's AI strategy and serves as the central hub driving synergy and efficient coordination.

During Chinese New Year, we launched our latest generation large model, Qwen3.5-Plus, which demonstrated outstanding performance across comprehensive benchmarks in reasoning, coding, and agentic capabilities. Through foundational architectural innovation, Qwen3.5-Plus significantly improved inference efficiency while reducing inference costs. Building on Qwen3.5, we will soon release our next version optimized for coding and agentic use cases.

On the consumer application side, the Qwen app, built on the strength of our models, has surpassed 300 million monthly active users. During Chinese New Year, we deepened its integration across Alibaba's ecosystem — connecting Qwen with Taobao Instant Commerce, Alipay, Fliggy, Damai, and Amap — giving it unique capabilities across the lifestyle domain and making it China's first all-in-one personal AI assistant for life, work, and learning.

We also recently launched Wukong, our AI-to-B AI agent platform. Wukong is the world's first AI-native enterprise-grade agent platform. It enables AI-powered upgrades to enterprise workflows while remaining compatible with each organization's data permissions and management processes. It serves as the unified interface for Alibaba's AI capabilities in enterprise work environments. B2B commercial capabilities across Alibaba's entire ecosystem will be progressively integrated via technical formats like Wukong Skills, to support Wukong in becoming the ultimate AI work assistant.

On Alibaba's other strategic priority, the broader consumption sector, we continue to advance our strategic initiatives. This quarter, our quick commerce business further expanded in scale, with continued share growth, improved customer retention, and sequential improvement in both unit economics and average order value. At the same time, quick commerce and e-commerce demonstrated clear synergies, driving Taobao's monthly active consumers (MAC) to double-digit year-over-year growth.

That concludes my remarks. I'll now hand over to Toby to share the financial update.

Financial Highlights

Toby Xu

CFO, Alibaba Group

(Original)

Thank you, Eddie.

Our strategic priorities are clear: we remain focused on AI + Cloud and Consumption businesses. We are seeing great momentum with gains in technology, customer adoption, market share and user engagement.

On AI + Cloud, we have the full-stack AI capabilities with all three core elements - model, cloud infrastructure and chips - and leadership in each with Qwen, Alibaba Cloud and T-Head. We also operate the most comprehensive consumer ecosystem in China that can monetize through AI. The launch of Qwen App was a major milestone and it can bring our consumer applications together.

On consumption, our quick commerce business continued to gain GMV market share in December quarter, while unit economics and AOV also continued to improve.

Now let’s look at the financial results. On a consolidated basis:

Total revenue was RMB284.8 billion. Excluding revenue from Sun Art and Intime, revenue on a like-for-like basis would have grown by 9%.

Total adjusted EBITA decreased by 57%, primarily due to our strategic investments in technology-related innovation initiatives and consumption front including quick commerce business, partly offset by the improved operating results in Cloud business, and enhanced operating efficiencies across various businesses.

Our GAAP net income was RMB15.6 billion, a decrease of 66%.

Operating cash flow was an inflow of RMB36 billion.

Free cash flow was RMB11.3 billion, a decrease of RMB27.7 billion from the same quarter last year. We are re-investing our cash flow to be a leader in AI and quick commerce.

As of December 31, 2025, we held US$42.5 billion in net cash. Excluding debt with maturities beyond five years, our net position stands beyond US$60 billion. This balance sheet strength gives us confidence to re-invest for long-term growth.

Now let’s look at our consumption businesses.

Revenue from Alibaba China E-commerce Group was RMB159.3 billion, an increase of 6%.

Customer management revenue increased by 1%. The slow-down in revenue growth was primarily due to weaker transaction activities and phase-out of the impact of software service fee implementation.

The Taobao app achieved a double-digit increase in MAC during the quarter, driven by the growing mindshare and increasing scale of our quick commerce business.

Revenue from our quick commerce business increased 56% to RMB20.8 billion. During the quarter, we executed our plan to further grow the scale of our quick commerce business, improve user experience, improve UE and increase AOV month-over-month during the quarter.

Alibaba China E-commerce Group adjusted EBITA was RMB34.6 billion, a decrease of 43%, primarily due to the investment in quick commerce, user experiences, and technology.

Going forward, this adjusted EBITA will continue to fluctuate quarter-over-quarter due to intense competition and significant investment in user experience.

Revenue from AIDC grew 4% this quarter. AIDC’s adjusted EBITA loss narrowed significantly year-over-year, driven by a combination of logistics optimization and investment efficiency enhancement. The UE of the AliExpress’ *Choice* business also improved on a sequential basis.

Next, let’s look at the business updates and results of Cloud Intelligence Group.

Our Cloud business delivered another quarter of accelerating growth. Revenue from external customers grew 35%, up from 29% last quarter.

AI-related products continue to lead this momentum. We delivered our tenth consecutive quarter of triple-digit growth in AI revenue. Its share of external cloud revenue continued to increase. This is a clear reflection of the scale and acceleration in our AI business.

The adjusted EBITA margin remained relatively stable at 9%.

We will continue to invest in customer growth and technology innovation to increase adoption of AI cloud infrastructure and strengthen our market leadership.

All Others segment revenue decreased by 25% to RMB67.3 billion, mainly due to the disposal of Sun Art and Intime businesses, as well as the decrease in revenue from Cainiao, partly offset by the increase in revenue from Freshippo and Alibaba Health.

All Others adjusted EBITA was a loss of RMB9.8 billion, primarily due to the increased investment in technology businesses including Qwen models and consumer facing Qwen, partly offset by the improved results of Cainiao, Hujing DME and other businesses.

Qwen Models have become one of the most widely adopted open-source model families globally, surpassing 1 billion cumulative downloads on Hugging Face by the end of this January. And consumer-facing Qwen has surpassed 300 million MAU across platforms, which reinforces user engagement and expands long-term monetization potential.

We have been increasing investments on these technology fronts including the Spring Festival campaign. Building on the strong momentum as Eddie mentioned earlier, we will continue to invest substantially in Qwen models and Qwen app.

Our unallocated adjusted EBITA was a loss of RMB2.7 billion, compared to a loss of RMB0.2 billion in the same quarter last year, which reflected cost associated with talent retention incentive from the one-off replacement awards plan of Ele.me.

Thank you. We will now open up for Q&A.

Q&A

Lydia Liu:

(Original)

Hi, everyone. For today's call, you are welcome to ask questions in Chinese or English. A third-party translator will provide consecutive interpretation for the Q&A session. Please note that the translation is for convenience purpose only. In the case of any discrepancy, our management statement in the original language will prevail. If you are unable to hear the Chinese translation, bilingual transcripts of this call will be available on our website within one week after the end of the meeting.

Robin Zhu (Bernstein):

(Original)

Thanks, management, for taking my question.

Could you give us some specific examples of how Alibaba Token Hub will change how the different cloud and AI businesses work together going forward from an organizational standpoint? And strategically, what changes or goals are you hoping to achieve with this new structure that improves on the previous arrangement going forward?

And then if management could share a hierarchy of priorities in cloud and AI, is it market share and revenue growth, such as the target you just announced, versus having the best first-party model capabilities, versus consumer-side traction with customers using agentic AI, or anything else?

Eddie Wu:

(Translation)

Thank you for your question.

We established the Alibaba Token Hub (ATH) Business Group in response to the new era of technological transformation. From the second half of 2025 through the first few months of 2026, we have entered a new era driven by agents. The key difference between this agentic era and earlier phases of AI is the much closer integration between models and applications, which has become critically important for both model and application development.

In the early stages of AI, most pre-training relied on static datasets. Now, in the agentic era, many advances will stem from the close integration between models and business applications, as well as from data feedback loops generated in real customer usage scenarios.

Looking across the layers of AI — from the application layer to the model layer, down through AI infrastructure such as cloud computing, and then the chip layer — we believe that in this agentic era, the integration of models and applications has become a more critical factor.

Next, let me explain how the different businesses within ATH connect.

Based on the industry trends we are seeing, the future could not be clearer: the AI agent application layer is poised for massive ecosystem innovation and a proliferation of rich new applications. At this layer, our portfolio includes the Qwen app as our consumer-facing personal assistant, and Wukong, which we aim to build into the leading enterprise assistant for Chinese businesses. We also expect numerous industry-specific, vertical, and scenario-based AI agent applications to emerge, supporting growth across diverse sectors.

Hence, beyond the application layer, we need a very robust MaaS offering. MaaS serves as the channel linking the model and application layers. In addition to supporting our own internal applications, a robust MaaS service will enable a broad range of external AI industry applications. We see enormous market value and opportunity in this area.

So going forward, we see the AI application layer as the main channel through which AI tokens will be distributed. Stronger model capabilities will naturally attract more applications, and a more efficient, more capable MaaS product will bridge these two layers more effectively. This is the business logic underlying the new ATH structure.

In terms of the model and application layers, our top priority, absolutely, is to develop the most powerful and intelligent models. Only with the strongest models can we drive the expansion of application scenarios across industries and attract diverse industry applications to adopt our MaaS services.

However, I want to emphasize: in order to build the most powerful models, we must ensure close integration with our own consumer-facing and enterprise-facing applications, and with a host of diverse industry applications connected through MaaS. We need more users to engage with our models in order to create a data flywheel and a closed-loop business model. With more scenarios, more data, and more users, we can continuously iterate and enhance our model capabilities. Driving this flywheel effect is also part of the rationale behind the establishment of ATH Business Group.

To summarize, I would say our priority is to strengthen model capabilities. However, building stronger models requires joint efforts across the entire ATH organization — on the model side, the application side, and the MaaS side — in order to drive sustained enhancement of model capabilities over time.

Joyce Ju (Bank of America):

(Original)

Good evening, management. Congratulations on the solid progress you've made in Cloud and AI, and thanks for taking my question.

My question is we see CMR growth slowing notably in the December quarter, given the macro pressures. We have seen China's online retail sales only up 2% year-over-year in the fourth quarter 2025. But more recently, NBS data point to a re-acceleration in January and February.

Could you share your latest view on the CMR trends heading into the March quarter? And whether you have started to see any improvement in consumer sentiment?

Jiang Fan:

(Translation)

Let me take this question.

As you mentioned, growth in the December quarter faced considerable challenges due to weak macro consumption, a warm winter, and the later timing of this year's Chinese New Year. Furthermore, because of the longer promotional window, our spending on consumer incentives was higher than in past years. As a result, CMR and e-commerce operating results growth softened in the December quarter.

However, since the beginning of the March quarter we have seen a very clear recovery in consumption. Propelled by this broader recovery and the added momentum from quick commerce transactions, our physical goods GMV and CMR growth have rebounded significantly, with marked improvement in e-commerce operating results.

Gary Yu (Morgan Stanley):

(Original)

Thank you, management, for taking my question.

My question is related to quick commerce. I understand that in the past couple of months, we have achieved certain milestones in terms of GMV market share and also UE improvement. How should we look at the priority going forward? Are we aiming for market share or hoping to take this opportunity to improve unit economics, reduce loss?

And how should we look at the synergy between quick commerce and traditional e-commerce? And how should we see these synergies to translate into CMR better growth going forward?

Jiang Fan:

(Translation)

Let me take this question.

While growing our market share, we have continued to improve our UE, driven by heightened logistics efficiency, stronger monetization capabilities, and a more optimized order mix. We expect to further optimize UE in the coming quarters.

Over the past year, quick commerce has clearly driven growth across our platform. Our broader e-commerce business, including quick commerce, added 150 million AACs, including 100 million conventional e-commerce physical goods AACs, which is more than the previous three years combined.

In the short term, new consumers exhibit lower average order values and purchase frequencies than existing users'. We aim to continuously increase their ARPU and frequency, which will serve as a new growth engine for our platform. At the same time, quick commerce is clearly driving sales in relevant categories — particularly food, fresh produce, and healthcare — while contributing to contributing to Freshippo and Tmall Supermarket's accelerated growth.

Looking ahead, we remain committed to our target of achieving over RMB 1 trillion GMV in quick commerce by FY2028. Once that GMV target has been achieved, we expect quick commerce to generate positive cash flow at scale, and to be profitable in FY2029.

Quick commerce has become a cornerstone of our e-commerce business, playing a strategically vital role for Taobao and Tmall's e-commerce businesses and in the AI era by driving user acquisition and engagement, fulfilling diverse consumer needs, increasing transactions and enhancing monetization, and supporting logistics infrastructure. We are committed to investing decisively in quick commerce over the next two years in order to achieve the RMB 1 trillion GMV target and secure market leadership.

Alicia Yap (Citigroup):

(Original)

Good evening, management. Thanks for taking my question.

I have some questions regarding your chip business, T-Head. So, there have been reports that Alibaba plans to spin off the T-Head unit as a separate listing. Can management provide any information about this? And if so, what is the expected time frame for this to occur?

And in the meantime, can you share more operating metrics? So, in addition to the 470,000 chips that you mentioned you shipped to external customers, how we reconcile that number, the shipments to the revenue side? And also, what is the expected growth rate for your chip business in the coming year?

And I think you mentioned that currently 60% of these chips are serving external customers. So maybe can you also share with us, are these chips for external customers mainly used for inferencing? And then for internal, is it used for model training and also inferencing?

And then lastly, how do the T-Head chips compare to other domestic chips? If management can share some details, it would be great.

Eddie Wu:

(Translation)

Thank you for the question. T-Head is an important component of Alibaba's full-stack AI strategy, and I'd like to take this opportunity to elaborate.

Within China's AI chip ecosystem, we believe that T-Head is in the top tier in terms of technology capabilities and product capabilities. Our products cover the entire AI workflow, from training and fine-tuning to inference. T-Head AI chips are already deployed at scale via Alibaba Cloud, both for training workloads and for Model Studio (Bailian) inference workloads.

Over 60% of T-Head chips are being used by external commercial customers across Alibaba Cloud's public cloud and hybrid cloud offerings. These customers span the internet, finance, autonomous driving, and smart manufacturing sectors, and they are using T-Head chips for both training and inference. Moreover, our software stack is highly compatible with the CUDA ecosystem, enabling customers to migrate their systems with minimal time and effort.

Another point I would make relates to T-Head's significance for Alibaba. Because domestic chips still lag behind their foreign counterparts in manufacturing process nodes and overall performance, we aim to deepen our co-design with Alibaba Cloud infrastructure and our Qwen models to deliver superior price-performance. This is what differentiates us from other chip companies. Our goal is to create AI capabilities that offer superior value for money. T-Head is a key driver in lowering inference costs on the Model Studio platform.

An additional point is that, beyond improving overall AI efficiency and reducing costs, given the unique circumstances of China's AI industry today, T-Head provides major value to us by guaranteeing the supply of AI computing power. I believe that over the next three to five years, global AI compute capacity will be in extremely short supply, and this shortage will be even more acute in the Chinese market. As the only cloud provider in the Chinese market with in-house chip design capabilities, T-Head is crucial to Alibaba Group. Providing a greater supply of AI compute will help our cloud and AI businesses, including our MaaS offerings, achieve stronger growth momentum.

Finally, let me share some of our expectations for the future.

Over the past two years, we have successfully shipped over 470,000 T-Head chips, which are being deployed across various commercial scenarios ~~deployed~~, with annualized revenue reaching the 10 billion RMB level. In 2026 and 2027, we expect T-Head to continue scaling the production of high-quality AI chips. This will provide robust computing power to support the Group's AI business, serving as a powerful growth driver for our overall AI initiatives while also significantly boosting our future profitability.

In summary, T-Head's value to Alibaba goes beyond cost optimization. It primarily serves to ensure supply chain resilience. In an era of scarce computing power, I see this as crucial to Alibaba's AI strategy.

While we do not rule out a potential IPO for T-Head in the future, we currently do not have any specific timetable.

Yuan Liao (CITIC Securities):

(Translation)

Good evening, management. Thank you for taking my question.

My question is about the AI revenue target you just shared. Could management provide more color on the goal to exceed US$100 billion in revenue in the next five years? For example, if this period runs through 2031, what is the implied CAGR? Also, could you break down the key growth drivers and how we should think about them?

Finally, given this revenue growth at scale, when can we expect to see Alibaba Cloud's margins trending upward?

Eddie Wu:

(Translation)

Thank you for your question.

We have fairly strong visibility toward achieving our goal of exceeding US$100 billion in AI and cloud revenue over the next five years, given the market growth we are seeing and our robust product foundation.

I believe our biggest growth driver will come from continued breakthroughs in foundation model capabilities. In the first few months of 2026, we've seen a clear trend: these models are now capable of executing complex enterprise workflows. As more companies deploy AI agents to handle end-to-end tasks, the market that AI and cloud services have traditionally targeted within enterprise IT budgets is undergoing a fundamental shift. When enterprises consume tokens at scale, they no longer see it as IT spending. Instead, they view it as part of their production or R&D costs, a core operational input. This is the most fundamental underlying factor driving our future AI growth.

We see the biggest growth momentum coming from three areas:

First, MaaS, driven by large models, will be our core growth engine. MaaS supports our own applications, our clients' applications, and a wide range of AI use cases across industries, including AI software. Growth driven by MaaS will be a key contributor to future AI and cloud revenue.

Second, there is a very significant incremental market for enterprise-level internal inference and training. While public MaaS will be a massive market, many mid-to-large enterprises will continue to require internal inference and training. This market will persist over the long term because enterprises will choose private deployments for certain applications based on their specific business models, security requirements, or unique use cases. Some scenarios will be well suited to public MaaS API services, while others will require private deployment. These private deployment scenarios represent a major incremental opportunity for Alibaba Cloud's AI infrastructure.

Third, a significant but often overlooked growth driver is the expansion of traditional, CPU-centric cloud computing in the era of AI agents. Traditional cloud computing was designed for human IT engineers. In China, there are perhaps a few million—at most 10 million—engineers who are potential customers. But as large models spawn billions of AI agents, these agents will require massive runtime environments supported by traditional CPU-centric cloud infrastructure. They will require CPUs, databases, storage, and large amounts of memory to support their long-term operation.

As a result, transforming traditional cloud computing from products designed for human engineers into products optimized for AI agents presents tremendous growth potential. Executing this transformation is a key upgrade focus for Alibaba Cloud this year.

As our revenue scale continues to grow, our AI business is undergoing a major business model upgrade: shifting from selling computing resources to selling intelligent capabilities. Combined with the cost reductions and efficiency gains from our proprietary T-Head chips, we expect cloud profit margins to improve clearly and sustainably as AI and cloud revenues scale. However, this improvement will not be linear. Instead, it will likely occur in step changes, driven by sudden shifts in economies of scale, massive scaling of our T-Head chips, and the broader realization of scale effects across our products.

Regarding the 2026 to 2031 CAGR you asked about: you can plug the numbers into your calculator and calculate the implied CAGR. But in practice, our R&D investments and market growth will not unfold linearly. Investments we make today may take one or two years to generate significant growth. But overall, we are highly confident in achieving our five-year revenue target.

Alex Yao (JP Morgan):

(Translation)

Thank you, management, for taking my question.

I'd like to shift gears and look at the current development phase of our e-commerce business. You previously mentioned that e-commerce would enter a three-year investment cycle. I'm wondering if that three-year timeline is now being adjusted given the new opportunities we are seeing in quick commerce, food delivery, and agentic commerce?

If there are no adjustments, does that mean we are currently around the midpoint of this three-year cycle? Going forward, should we expect to see gradual financial improvements, leading into a relatively stable financial harvest period by the end of year three?

Could you share how you're thinking about our current positioning and overall direction in e-commerce?

Jiang Fan:

(Translation)

Let me take this question.

As I mentioned earlier, we are making significant investments in quick commerce this year, as we see a highly definitive opportunity. Over the next two years, we will continue to invest toward our goal of surpassing RMB 1 trillion in quick commerce scale. We believe that in two years' time, these investments will deliver positive economic returns for our e-commerce business as a whole.

Let me also touch on AI. Eddie has already spoken at length about its impact, and I believe that AI will profoundly reshape e-commerce as well. But a three-year horizon is an extremely long timeframe in AI, given how quickly it evolves week by week and month by month. That is why we are actively investing in AI initiatives. This year, we will continue to roll out new AI-driven experiences for both consumers and for merchants, and we will help merchants upgrade their operating models with AI.

The emergence of AI will also drive massive upgrades across many e-commerce sub-sectors. For example, it will create tremendous opportunities in our B2B business and similar segments, and we intend to actively seize these new opportunities.

Lydia Liu:

(Original)

Okay. That wraps up the Q&A session of today's earnings call. Thank you very much for joining us today, and we look forward to speaking with you soon.

[END OF TRANSCRIPT]