



Alibaba Group

March Quarter 2026 and Full Fiscal Year 2026 Results

Conference Call

Transcript

Wednesday, 13 May 2026

Introduction

Lydia Liu

Head of Investor Relations, Alibaba Group

(Original)

Good day, everyone. Thank you for joining Alibaba Group's March Quarter and Full Fiscal Year 2026 earnings call.

On the call with me are Joe Tsai, Chairman; Eddie Wu, Chief Executive Officer; Toby Xu, Chief Financial Officer; Jiang Fan, Chief Executive Officer of Alibaba E-commerce Business Group.

I would like to remind you that this call is also being webcast on our corporate website. A replay of the call will be available on our website later today.

Now I will quickly cover the safe harbor. Today's discussions may contain forward-looking statements, particularly statements about our business prospects and expected financial results, that are subject to risks and uncertainties, which could cause actual results to differ materially from those contained in the forward-looking statements. Please refer to the safe harbor statements that appear in our press release and investor presentation provided today. Please note that certain financial measures that we use on this call are expressed on a non-GAAP basis. Our GAAP results and reconciliations of GAAP to non-GAAP measures can be found in our earnings press release and investor presentation provided today.

Unless otherwise stated, growth rate of all stated metrics mentioned during this call refers to year-over-year growth versus the same quarter last year.

With that, let me turn the call over to Eddie.

Group Business Update

Eddie Wu

CEO, Alibaba Group

(Translation)

Welcome to Alibaba Group's fiscal year 2026 fourth quarter earnings call.

Over the past quarter, Alibaba's high-intensity investment in our two strategic priorities of AI + Cloud and consumption is rapidly translating into tangible business results, with Group revenue growing 11% year-over-year¹.

This quarter, Cloud Intelligence Group's external revenue growth accelerated to 40%, and AI-related product revenue achieved triple-digit growth for the eleventh consecutive quarter. China e-commerce CMR grew 8% year-over-year on a like-for-like basis, and the quick commerce business achieved significant UE (unit economics) improvement while maintaining market share.

We are at a pivotal inflection point in the evolution from conversational chatbots to autonomous AI agents, which is directly driving explosive growth across three core workload categories: AI training, inference, and agent orchestration.

Against this backdrop, Alibaba's AI has moved beyond the initial investment phase and progressed commercialization at scale.

Let me now walk you through four areas in detail: AI commercialization, cloud infrastructure, the AI application ecosystem, and our consumption businesses.

1. The AI and cloud commercialization inflection point has arrived.

This quarter, Cloud Intelligence Group's annualized AI-related product revenue has surpassed RMB 35.8 billion, continuing to maintain triple-digit growth. AI-related product revenue now accounts for 30% of Cloud Intelligence Group's external revenue. We expect that in about one year, AI-related product revenue will cross the 50% threshold, becoming the primary engine driving the cloud business's revenue growth. As a result, Cloud Intelligence Group's external revenue growth is expected to continue accelerating beyond its current 40% rate over the coming quarters. Given the certainty of long-term AI demand and our full stack technology advantages, we expect this trajectory to sustain strong growth over the medium to long term.

¹ This refers to that excluding revenue from Sun Art and Intime, Group revenue would have grown by 11% year-over-year on a like-for-like basis.

This reflects AI's role in driving a comprehensive upgrade of Alibaba Cloud's entire business, as its growth engine fully pivots from traditional compute and storage to models, AI compute, and agent services.

We are also seeing exponential growth in AI Model and Application Services revenue, a new revenue engine driven jointly by foundation model services and AI-native software.

Over the past three months, token consumption volumes on our model services platform grew substantially quarter-over-quarter, as enterprise customers accelerated their shift from simple tasks to production-scale and complex workloads, driving continued growth in demand for model and application services on the Model Studio platform.

We expect model and application services annualized recurring revenue (ARR), inclusive of the Model Studio platform, to surpass RMB 10 billion in the June quarter and RMB 30 billion by year-end. The high-margin profile of this revenue stream is becoming increasingly apparent, making it a source of healthy, high-quality growth.

2. Our AI infrastructure underpins our full technology stack and constitutes a durable moat.

T-Head's proprietary GPU chips have achieved scaled mass production, with over 60% of compute capacity already serving external customers across Internet, financial services, and autonomous driving verticals. As the only AI cloud provider in China capable of delivering self-developed AI chips at scale, we have secured autonomy over our compute supply chain while providing customers with highly competitive AI inference and training services. In an environment of compute scarcity, this structural advantage is favorable to our revenue growth and gross margin improvement.

At the same time, our cloud products are accelerating their AI-oriented upgrade. The surge in agent workloads has significantly elevated demand for traditional cloud products built around CPU, storage, and containers, and we are upgrading these into infrastructure solutions optimized for the agent era.

3. At the application layer, we have built a complete closed loop spanning AI-native software to a full agent ecosystem.

The newly established Alibaba Token Hub (ATH) continues to launch new products, connecting consumer and enterprise environments, with breakthrough progress in AI-native software and coding agents.

Our Qwen model continues to iterate across reasoning, coding, and agentic capabilities.

On the enterprise side, we have launched a range of products spanning intelligent workplace tools, AI coding, and business operations management, helping enterprises unlock greater productivity.

On the consumer side, Qwen app fully integrated Taobao and Tmall's commerce service capabilities on May 7. With this, Qwen app is now deeply embedded across the Alibaba

ecosystem, spanning Taobao, Alipay, Amap, and Fliggy, making it China's first all-in-one personal assistant to seamlessly bridge everyday life, productivity, and learning.

4. Across our consumption businesses and at the group level, we are prioritizing long-term value.

Beyond AI, our consumption strategy continues to progress steadily, with CMR growth rebounding significantly. This quarter, CMR grew 8% year-over-year on a like-for-like basis, as we continue to improve user experience and merchant operating efficiency. The quick commerce business achieved significant unit economics improvement while maintaining stable market scale.

In summary, the return on our investments in AI + Cloud and consumption is increasingly clear: AI + Cloud revenue growth is accelerating with improving margins, model and application services ARR continues to grow at pace, and operating efficiency across our consumption business continues to improve. Facing the historical opportunity that AI represents, Alibaba is at a pivotal juncture where our technology investments are beginning to pay off commercially. We will maintain our strategic resolve and leverage our full-stack AI capabilities to support long-term growth.

That concludes my prepared remarks. I will now hand over to Toby to walk you through our financial results. Thank you.

Financial Highlights

Toby Xu

CFO, Alibaba Group

(Original)

Thank you, Eddie.

Our strategic priorities remain laser-focused on AI + Cloud and Consumption businesses. Multiple growth catalysts, including technology advancement and business innovation, are aligning to create strong tailwinds.

On AI + Cloud, our full-stack capabilities span models, cloud infrastructure and applications, with established leadership in every layer. The strong growth of our AI + Cloud businesses and clear path to monetization of our MaaS platform gives us confidence to make significant investments to extend our leadership.

On Consumption, we achieved strong CMR growth on a like-for-like basis during the quarter, and our quick commerce business continued to improve UE and AOV quarter-over-quarter.

Now let's look at the financial results for this quarter. On a consolidated basis:

Total revenue was RMB243.4 billion. Excluding revenue from Sun Art and Intime, revenue on a like-for-like basis would have grown by 11%.

Total adjusted EBITA decreased 84%, primarily due to our strategic investments in technology businesses, quick commerce and user experience, partly offset by the improved operating results supported by continued growth in customer management service and Cloud business, and enhanced operating efficiencies across various businesses.

Our GAAP net income was RMB23.5 billion, an increase of 96%, primarily attributable to the year-over-year increase in net gain from mark-to-market changes of our equity investments, and disposal losses of Sun Art and Intime in the same quarter last year, partly offset by the decrease in adjusted EBITA.

Operating cash flow was an inflow of RMB9.4 billion.

Free cash flow was an outflow of RMB17.3 billion. We are re-investing our cash flow to enhance our competitive advantage in AI.

As of March 31, 2026, we held approximately USD38 billion in net cash. Excluding debt with maturities beyond five years, our net cash position stands at approximately USD59 billion. This balance sheet strength gives us confidence to invest for growth.

Now let's first look at our consumption businesses.

Revenue for Alibaba China E-commerce Group was RMB122 billion, an increase of 6%.

Customer management revenue increased by 1%. To help merchants grow their businesses and increase willingness to spend on our platform, we upgraded our business development program for select merchants during the quarter, under which the level of platform subsidies for these merchants is directly tied to their marketing spend on our platform. For accounting purposes, such subsidies previously recorded as sales and marketing expenses are now recorded as a contra revenue item to CMR. Accordingly, CMR grew 1% year-over-year during the quarter. Excluding the contra revenue impact from the program, on a like-for-like basis, CMR would have grown 8% year-over-year.

Revenue from our quick commerce business increased 57% to RMB20 billion. The quick commerce business further improved UE, and increased AOV quarter-over-quarter primarily driven by order mix optimization.

Alibaba China E-commerce Group adjusted EBITA was RMB24 billion, a decrease of 40%, primarily due to the investment in quick commerce, user experience, and technology, while there is positive contribution from customer management service.

Excluding loss from our quick commerce business, our Alibaba China E-commerce Group EBITA would have been stable year-over-year and will fluctuate quarter-over-quarter due to significant investment in merchant retention and user experience.

Revenue from AIDC grew 6% this quarter.

AIDC's adjusted EBITA loss narrowed significantly year-over-year, approaching break-even, driven by a combination of logistics optimization and operating efficiency. The unit economics of the AliExpress' Choice business continued to improve substantially on a sequential basis.

Next, let's look at the business updates and results of Cloud Intelligence Group.

Our cloud business delivered another quarter of accelerating growth. Revenue from external customers accelerated to grow 40%.

AI-related products continue to lead this momentum. We delivered our eleventh consecutive quarter of triple-digit growth in AI revenue. Its share of external cloud revenue continued to increase, now accounting for 30%.

This quarter's AI revenue is RMB9 billion and the annual revenue run rate is approximately RMB36 billion (or US\$5.3billion)². This is a clear reflection of the scale and acceleration in our AI business.

The adjusted EBITA margin remained relatively stable at 9.1%.

² Annual revenue run rate is calculated using current quarter revenue multiplied by four.

All Others segment revenue decreased by 21% to RMB65.5 billion, mainly due to the disposal of Sun Art and Intime businesses, as well as the decrease in revenue from Cainiao, partly offset by the increase in revenue from Freshippo and Amap.

All Others adjusted EBITA was a loss of RMB21.2 billion, primarily due to the increased investment in technology businesses including foundation models and the consumer-facing Qwen app.

As we close this fiscal year, we remain committed to delivering consistent shareholder returns. Our board of directors has approved an annual dividend of US\$1.05 per ADS.

We will continue to invest decisively in AI and consumption businesses, where we see significant long-term growth potential and our competitive advantages are compounding. We believe these investments to deliver growth and returns over time, ultimately creating greater value for our shareholders.

Thank you. We will now open up for Q&A.

Q&A

Lydia Liu:

(Original)

Hi everyone, you are welcome to ask questions in Chinese or English. A third-party translator will provide consecutive interpretation. In the case of any discrepancy, our management's statement in the original language will prevail.

Ronald Keung (Goldman Sachs):

(Original):

Thanks for sharing the very sizable AI MaaS and applications ARR scale and the target for the first time.

I just want to ask, how much of that ARR is driven by our in-house models like Qwen versus third-party models? And given the recent token price hikes, what would be the implications to MaaS and also our Cloud margins as a result?

Eddie Wu:

(Translation)

This quarter, we disclosed a new metric: our model and application service revenue. This revenue currently comprises two main components: API service revenue from our Model Studio MaaS platform, and subscription revenue from our AI-native software.

At present, the large majority of this revenue comes from Model Studio MaaS API services. That said, Alibaba Cloud's Model Studio is a relatively open platform on which we host both our own proprietary models and third-party models, including both open-source models and closed-source models. In terms of revenue scale, the vast majority still comes from our proprietary models, including Qwen, our foundation models, as well as our voice and video models.

Your second question is also a key one. Over the past quarter and in recent months, the industry has undergone a very significant shift: AI is transitioning from conversational chatbots to agentic workloads. Agents help customers complete highly complex reasoning tasks, which has driven a sustained, steep rise in customer demand for model inference.

Because our models can now help customers complete far more complex work, customer acceptance has remained very high even after the increase in overall API token prices, and demand has been consistently strong. In fact, as things currently stand, our supply is still not able to fully meet customer demand, and we have a long waitlist of clients.

Against this backdrop, we believe that, first, the MaaS business inherently carries a higher gross margin profile compared to the IaaS business.

Several other points are also noteworthy. First, inference technology continues to advance, and each quarter we are seeing improvements translate into sustained, incremental gains in token output capacity per server and per GPU.

Second, as model capabilities continue to strengthen, we expect model pricing to remain on a sustained upward trajectory over the next one to two years. From this perspective, we believe the rapid growth of our MaaS business will have a significantly positive impact on our overall gross margin in the coming quarters.

Kenneth Fong (UBS):

(Original):

Congrats on the very strong progress on the AI.

I have a question regarding the return on invested capital on the AI investments. While our AI investments have driven impressive 40% Cloud external revenue growth, they have also created significant drag on the Group's free cash flow as well as our EBITA. So how should investors assess the return on this investment? What is the management framework for balancing the aggressive AI spending versus earnings stability?

Toby Xu:

(Translation)

Let me address the first part, as I imagine many investors will be interested in how we manage free cash flow, given that it was negative this quarter.

First, as you can see, the negative free cash flow was primarily driven by our AI investments over the past year. It is precisely because we have recognized this historic opportunity that we have been highly resolute in making these investments. This is the main reason for the negative free cash flow.

Looking ahead over the next two years, we will remain equally resolute in making these investments, because this window of opportunity may only last a few years for us.

At the same time, there has been no material change in our operating cash flow. Let me highlight a few points.

First, on the consumer side, Taobao and Tmall are our most important contributors of operating cash flow, and overall operating cash flow generation is very stable. Over the next two years, as Quick Commerce losses narrow significantly and AIDC moves toward profitability, we expect our overall consumer business to deliver very strong positive operating cash flow. That's the first point I want to make.

Second, also very importantly, our investments in Cloud infrastructure will continue to drive revenue acceleration and gross margin improvement in our AI + Cloud business, as Eddie discussed earlier. This will further enhance operating cash flow from our Cloud segment, which can also support our continued investment in cloud infrastructure.

Third, we have a very strong balance sheet. Our current net cash position is approximately US\$38 billion. If we exclude debt with maturities beyond five years, our net cash is approximately US\$59 billion. This balance sheet strength provides solid support for our investments in cloud infrastructure.

Finally, we also have strong financing capabilities in capital markets and can access various forms of financing to meet our strategic development needs.

So that's my response to your question on cash flow, including on how we are positioning ourselves to ensure ample cash flow.

Let me now turn it over to Eddie to see if he has anything to add.

Eddie Wu:

(Translation)

On your question about our AI investments and the future ROI, let me take this opportunity to share our views as well.

Looking at how the AI business is developing today, I think it is like a manufacturing business. In order to be able to generate more revenue, we need to build two core "factories". The scale of these two factories will determine how large our future revenue scale can be.

The first we can call the "AI training factory", and the second, the "AI inference factory". Both of these "factories" are underpinned by the build-out of AI data centers, which I believe will consume a significant amount of the Group's free cash flow.

However, the path to returns on our investment in this essential data center infrastructure is very clear. These data centers have a clear commercialization path on the enterprise side, including through Cloud IaaS services, our MaaS platform, and our AI-native applications which

generate additional revenue on top of the models. Today, there is virtually not a single idle GPU in our servers.

Given the demand we anticipate over the next three to five years, we have invested heavily in AI data center construction, and there is a high degree of certainty in the investment returns.

Thomas Chong (Jefferies):

(Original):

My question is on Quick Commerce.

You have talked about the improvement in UE in the prepared remarks. I just want to get some more color about the drivers behind in terms of AOV, subsidies ratio, fulfillment ratio, etc.

On top of that, I remember last time we talked about the outlook for Quick Commerce over the next couple of years. Is there any update or changes in terms of how we think about the landscape or the UE in the next three years?

Jiang Fan:

(Translation)

First, as a result of a year of investment, we have achieved rapid growth in our Quick Commerce business, marking a fundamental shift in our market share. Compared to the same March quarter last year — that is, before our large-scale investment—both our order volume and market share have increased significantly. Overall order volume in January-March was 2.7 times that of the same period last year, with non-food orders tripling year-over-year.

Since April, while maintaining order volume, we have continued to drive substantial improvement in UE through enhanced fulfillment logistics efficiency, order mix optimization, and further increases in AOV in particular. We are confident that UE will turn positive by the end of FY27.

Alongside UE optimization, we will continue to innovate to improve the experience for both consumers and merchants, thereby sustaining our long-term competitiveness in Quick Commerce. We are confident that our Quick Commerce business will achieve overall profitability in the future at the scale and market share.

This quarter, Quick Commerce continued to generate synergies with our conventional e-commerce business, particularly by driving new customer acquisition, enhancing user engagement, fulfilling diverse consumer needs, increasing transactions, improving monetization, and supporting logistics infrastructure. All this helped drive the overall development of Taobao and Tmall.

We continued to see Quick Commerce driving sales growth in relevant categories, especially food and fresh produce, contributing to Freshippo and Tmall Supermarket's accelerated growth.

Our conventional e-commerce business posted relatively strong growth in the March quarter in both GMV and CMR on a like-for-like basis. Quick Commerce played a definitively positive role in driving this performance.

Jialong Shi (Nomura):

(Translation)

I want to follow up on the MaaS business you discussed in your opening remarks. I'd like to get your thoughts on Alibaba's competitive advantages in the MaaS space compared to the other major AI platforms and AI start-ups in China.

In the United States, we've seen AI agents, and AI coding in particular, emerge as the fastest-growing area for AI commercialization. I am wondering when you expect we might see similar growth rates for AI coding in China?

Separately, Chinese enterprises have historically been less willing to pay for SaaS products. Do you think that makes the commercial prospects for AI coding products in China less promising than in the US?

Eddie Wu:

(Translation)

Alibaba Cloud's MaaS platform, Model Studio, is positioned as an open AI inference platform. Currently, the majority of its revenue is driven by our proprietary models. In contrast to the AI startups in China, we are investing at a much larger scale and across a much broader range of model types.

That said, because these startups tend to focus on a single model or a narrow vertical, they have strong technology and can move rapidly ahead with better strategic intuition. Purely in terms of the MaaS space, those AI startups are partners for us in a certain degree.

At Alibaba, however, we emphasize broad-based R&D across all model domains, from our coding-focused Qwen foundation models to our video models, including Wanxiang and HappyHorse, our more forward-looking world models, and our speech models. We believe that going forward, users across diverse scenarios will need to draw on multiple model capabilities in combination to meet their business needs.

So that's a key difference between us and the leading AI startups. And at the same time, these startups can also be cooperation partners for Model Studio to a certain extent.

Your question on AI Coding is also a good one. You asked when China might see similar growth rates, and our view is that China is already there. Based on what we're seeing on our own Model Studio platform, as well as what is being observed by others in the industry — including AI startups that cooperate with us — the surge in API demand from last November–December to May this year has been driven largely by improvements in AI Coding capabilities.

AI Coding isn't simply about replacing the work of software engineers. With the advancement of AI model capabilities, and with the integration of AI Coding into the broader agent runtime environment — what is known as “harness engineering” — along with the integration of related data domains, we are seeing AI Coding-driven, complex-task agents that can now complete tasks across virtually every type of digital work.

So, whether in the US or in China, the growth in this current wave of AI demand is being driven primarily by improvements in AI Coding. When these enhanced capabilities are integrated with computers and digital tool environments, in theory they can handle nearly all of the complex tasks of digital work going forward. This will be a very important growth trend over the next two to three years.

Regarding your comment about Chinese enterprises historically having lower willingness to pay for SaaS, I believe the large model era represents a fundamental shift. As AI models become increasingly powerful and genuinely capable of completing complex tasks in users' actual workflows, what we're seeing — across both US and Chinese enterprises — is that the demand to pay for AI capabilities that deliver real work outcomes is equally strong in both markets. In theory, as long as the value created by this work within the enterprise exceeds the token cost, demand for API tokens becomes, in a sense, effectively unlimited. From this perspective, we see the underlying growth in AI demand as having very strong long-term certainty. Let me also share some numbers. On the Model Studio platform, growth has been extremely rapid. We expect May-June data to show more than a tenfold increase from November-December levels. Recently, our ARR already exceeded RMB 8 billion, and I think it's highly certain that we can achieve an ARR of over RMB 10 billion this quarter.

Ultimately, what we're seeing is that the willingness to pay for AI intelligence that helps complete real-world business tasks is universal across Chinese and US enterprises alike.

Ellie Jiang (Macquarie):

(Original)

Just wanted to stay on the topics of global comparisons. If you look at globally, the overseas peers seem to have captured the most immediate ROIs in enterprise agentic workflows, whereas for the consumer and the monetization which remain a bit lagged. Going forward, considering that Alibaba is investing in multi fronts for infrastructure, models, Cloud, and Qwen app, how do we evaluate the strategic priority and resources allocation between 2B and 2C initiatives? If going forward enterprise side continues to gain more traction, will we consider gradually shift more resources away from Qwen app to Cloud and MaaS?

Eddie Wu:

(Translation)

Fundamentally, AI represents a revolution in computing paradigms. Ultimately, the goal of this revolution is to help users complete tasks or solve problems more effectively. In this regard, I see the enterprise and consumer sides as essentially the same.

That said, as things stand today — both globally and in China — the highest willingness to pay is indeed on the enterprise side, because ROI is easier to quantify. So the majority of our inference resources are currently allocated to enterprise-side commercialization.

But AI is ultimately positioned as an assistant to humans. Whether it functions as a work assistant, a personal assistant, or a learning assistant, the core objective is basically the same: AI is there to help people complete tasks. Those tasks may fall on the enterprise side or the consumer side, or span learning, work, and daily life.

We believe that on the consumer side, given current levels of user adoption and willingness to pay, a certain investment cycle will be required. But we also believe consumer AI assistants will also gradually develop their own business models as the technology advances, as user adoption grows, and as they are able to help users complete more and more tasks. We are already seeing this kind of business model in overseas markets. We expect to see more progress on consumer-side AI monetization in the domestic market within the next one to two years.

Joyce Ju (Bank of America):

(Translation)

I'd like to follow up on the earlier discussion about the future growth of the Cloud business. Beyond the topline acceleration, could management speak to the outlook for EBITA margin? Specifically, over the next few quarters as growth continues to accelerate, can we expect to see a margin expansion trend similar to that of overseas peers?

Eddie Wu:

(Translation)

We believe AI technology and its penetration across industries is still at an early stage. The primary objective for Alibaba Cloud and our AI business currently is growth in users, token consumption, and market share.

We aim to maintain growth that is faster than the market average in order to gain larger market share and firmly cement our absolute market leadership position. Those are the primary objectives, and margin is still secondary.

That said, there are several very clear industry dynamics at play right now.

First, it remains difficult to assess how AI industry demand will be met over the next three to five years, because of numerous physical bottlenecks — the construction cycle for AI data centers, the production cycles for chips and memory across the industry, the physical build-out of fabs, and overall capacity growth. Within a three- to five-year horizon, supply will struggle to keep up with the growth in AI demand.

Historically, IT CapEx has benefited from strong customer and asset scale effects. But because of the current tightness in supply, the cost of deploying a new server today is more than double what it cost us to deploy essentially the same hardware two years ago. Cost inflation has been over 100%.

This rising replacement cost for new servers is driving upward pricing power on cloud services for both existing and new customers. We believe this will have a positive effect on driving pricing power upward for all kinds of cloud services over the longer term.

Second, we're also seeing rapid growth in Alibaba Cloud's MaaS business. Because MaaS revenue is inherently a higher gross-margin business than our traditional IaaS and IT-asset-based businesses. At the same time, as inference technology continues to improve, per-GPU capacity keeps increasing. So what we're seeing now is that the same server, if allocated to Model Studio, generates higher revenue and gross margin than if it were allocated to traditional cloud computing services. As MaaS's revenue contribution grows, we believe it will lift our gross margin going forward.

Third, we have a full-stack technology advantage. T-Head's proprietary AI chips development and future production at scale will enable us to be potentially the most cost-effective inference platform in China. We believe this inference platform can generate superior gross margins and create strong synergies with our models.

In conclusion, because of these objective factors, we expect Alibaba Cloud's gross margin to improve meaningfully over the next one to two years. This momentum should be visible within the next one or two quarters.

Gary Yu (Morgan Stanley):

(Original)

I have a question regarding CapEx. What kind of level of CapEx investment is required in order to satisfy the demand from both MaaS and also the long-term Cloud revenue?

Also, management mentioned about T-Head opportunity. What is the current penetration of T-Head being deployed on Alibaba Cloud? As this penetration increases, what margin uplift we should expect from our in-house chip?

Eddie Wu:

(Translation)

Your first question is also an important one. In fact, during our prepared remarks last quarter, we shared the ambitious revenue target we have set for the next five years. That target represents a roughly 10x increase compared to Alibaba Cloud's external revenue in 2022 and 2023, prior to the AI large model boom.

By way of rough calculation, then, our data center footprint needs to expand at least 10x relative to 2022 in order to support our long-term business objectives. That is our overall target for our build-out, encompassing both our CapEx investments and the substantial compute we are currently securing through OpEx arrangements.

That means that the capital we need to deploy over the next five years to secure this compute capacity will far exceed the original budget of RMB 380 billion over three years that we had previously shared.

Of course, the situation is more complex now, and not all of these compute centers will be self-built using our own CapEx. For example, we will secure some through OpEx leasing arrangements. Also, as T-Head's chip production scales, we may sell T-Head AI servers to various compute or data center service providers and co-build with them.

So we will expand our data center capacity in a variety of different ways. But in terms of fundamental data center requirements, we are benchmarking against a scale that is 10 times our 2022 footprint, before the AI large model boom.

Regarding proprietary chip deployment on Alibaba Cloud, penetration remains at a relatively low level. To share further context, I'd like to highlight that T-head's chip portfolio extends beyond our own GPUs to include proprietary CPUs, storage chips, and networking chips. So going forward, we have the opportunity to run our full stack — GPUs, CPUs, storage and networking chips — all on proprietary silicon.

For now, penetration remains low due to various objective factors, including relatively limited domestic semiconductor manufacturing capacity, although it has been expanding in recent years.

As T-Head's full-stack proprietary chips gain penetration, the impact on our gross margins will be very substantial. That said, the magnitude of this impact is currently being moderated by several interacting factors. For example, because China's domestic semiconductor manufacturing process nodes still lag overseas ones, a chip of comparable specifications will still show some gap, whether in power consumption or efficiency, versus advanced foreign chips.

That said, gross margins on mainstream foreign AI chips are very high, basically in the 60% to 80% range. So even if domestic chips still fall somewhat short on overall performance or power efficiency, the high margins on advanced foreign chips leave domestic chips with enormous headroom for cost-performance gains.

How much of that cost-performance upside ultimately flows through to our overall gross margin will depend, as capacity expands, on the substitution ratio that domestic chips can achieve.

Lydia Liu:

(Original)

Thank you. This brings us to the end of today's earnings call. We appreciate your time and participation, and we look forward to speaking with you soon.

[END OF TRANSCRIPT]